



CO429

Cédric DESSEZ, Valentin GRAND, Yann NICOLAS,
Cyril NOVEL, Nicolas SERVEL



Parallel Algorithms Coursework

Cédric DESSEZ, Valentin GRAND, Yann NICOLAS, Cyril NOVEL, Nicolas SERVEL

November 29, 2013

Contents

1	Gathering the clues	2
2	Implementation	3
2.1	Sequential algorithm	4
2.2	Parallel MPI algorithm	4
2.2.1	Parallel version of the algorithm	5
2.2.2	Chunked method	6
2.3	Addition of OpenMP parallelization	7
2.4	Performance analysis	7
3	Conclusion	8
	More about the program	9
	Launching the program	9
	Overall architecture	9
	Problems interfacing MPI with the rest of the program	10
	BT telephone bill	11

1 Gathering the clues

The first thing we did was to analyze in depth the evidence given by our intel. Hopefully the engineers from Nukehavistan were kind enough to leave clues about the computer system.

The very first thing we did was to find the identity of the man portrayed on the system. We were able to process the face into the database of Interpol, hoping we will find a rogue scientist working for the Nukehavistan. Unfortunately no match were found. We then used the ReverseImages tool Q made for us and found that the man was *Leonhard Euler*, portrayed on a 1957 stamp of the USSR. Such a link supposes that the Nukehavistan is a communist nation and so a major threat to the democracy all around the world.

Because of the link with the USSR, we assumed that the funny writing above the clock was in Cyrillic script. As no one in our team was able to understand russian, we contacted the Translation Service of the MI-6. The translation was *fast inverter integrated functions*.

The text below the clock is in french. As we are french, we perfectly understood it and it means *Made in France*. However, it is very unlikely this machine was produced in France. The final design is ugly and everyone knows French are all about beauty and class. We assumed this was only a trick to harden our research.




Figure 1: The model of the classy French spy

We then took a look to the two sets of coordinates. The set along the board points to the lovely village of *Beaumont-en-Auge* in France. Among the famous people born in this town, one stands out. *Pierre Simon de Laplace* was a mathematician and his work could have highly interested the Nukehavistan. The second coordinates confirmed our idea : it points near the *Laplace* RER station in Paris. We don't believe in coincidence. Laplace's work is the key of the system. The power to the minus one on the second set of coordinates strengthens our assumption about the *inverter*.

We finally called the technical hotline. Since it is a charged number, we joined the bill at the end of the report. The technical hotline is only an answerbot. We tested if it was telling the truth by asking it the answer to the Ultimate Question of Life, the Universe and Everything. Since it answered 42, we assumed the bot was powerful enough to answer our upcoming questions¹. When we asked what was in the box, it said *Numerical Inversion of Laplace Transforms of Probability Distributions*. A quick and efficient DuckDuckGo search – we made sure we weren't tracked – pointed us to a research paper from 1995, presenting an algorithm using the Euler method for inverting Laplace transforms.

¹This raises serious issues about the advanced level of Nukehavistan in the Computer Science field. Maybe we should apply for a PhD in their university.

The loop was almost complete. We knew what the algorithm was and what to code. The only thing we didn't know was what the program was actually doing. We tried to reverse engineer the *mystery.o* file. We obtained some strange figures.

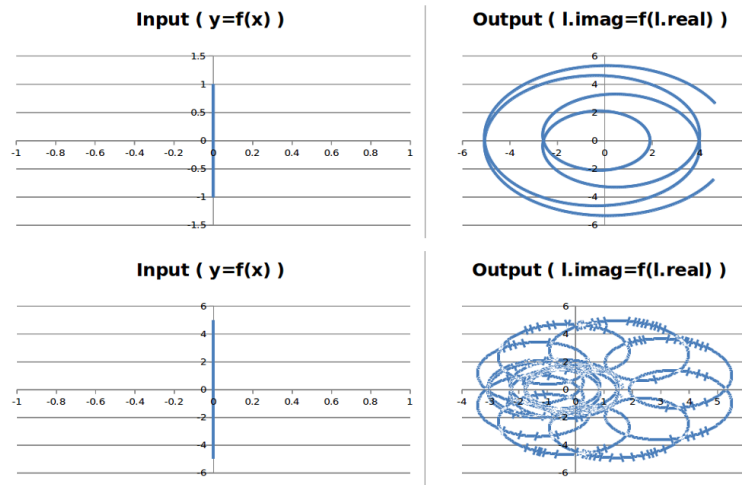


Figure 2: Reverse engineering inputs and outputs.

We then made two assumptions. The first possibility is that the Nukehavistan is trying to communicate with superior life forces. The second possibility is that this computer system is prime in the job of the Nukehavistan spies. Since we are fierce scientists, we ruled out the first option. Thus we brainstormed over the second option by watching all the James Bond movies. At the end of this intensive marathon, we reached a conclusion :



Figure 3: Watches

Every spy needs a good watch! If we look closely to the computer system, the small peak of the output is on the 3 and the big peak is on the 11. The clock has its small hand pointing between 3 and 4 – so in the 3rd hour. The large hand is pointing at 11. So the output of the computer system gives us the time!

2 Implementation

Since every good spy must be able to adapt and exploit their resources in the best possible way, they need a piece of software that is the fastest possible not only for an on-stage operation,

but also when they are preparing the mission at the headquarters with access to the best super calculator in the world. This is why we have optimized this tremendous clock not only to run sequentially on a single machine, but also in parallel over several processes, multithreaded or not (depending on the given option), and distributed among a pool of machines.

2.1 Sequential algorithm

The first step in building this powerful clock is to implement and optimize the sequential algorithm. We first made a rough copy of the algorithm from the research paper and then modified it to make it more efficient and more flexible.

In the paper, the parameters `N` and `M` are hardcoded respectively to the values 15 and 11. That allows them to have a quite efficient processing and to hardcode the binomial coefficients. Those parameters determine the accuracy of the approximations done by the algorithm. After running a few tests, we realized that the more influent of those is `N` which we decided to set by default to 100 (this also gives more computation to parallelize in the next steps). We have increased `M` a bit, up to 15, to improve the accuracy, but also in anticipation of the parallelization which will use 16 machines.

In order to plot the inverted Laplace transform, we need to run this algorithm many times to obtain enough points. Thus we have strived to mutualise as much computation as possible. First the binomial coefficients are processed once and for all and saved to memory, as well as their sum which is in fact 2^M (hardcoded to 2048 in the paper). Likewise, we process $\exp(\frac{A}{2})$ only once for the whole set of points.

Besides, we have little interest here in the truncation error estimate, so we have got rid of the part of the code that processed it in the paper.

We have also merged the last two loops. As a matter of fact, filling an array `SU` with one loop to get `AvgSu` by summing its elements in a second loop is unefficient because we can do both at the same time which reduces the loop administration overhead and allows us to keep only a `double` variable instead of the array `SU`.

We have also optimized the calculus of the variables `i` and `Y`. In the paper they are computed with multiplications of float number, whereas their evolution corresponds to an incrementation by a fixed value at each iteration. Since in general an addition costs much less than a multiplication, this makes the loop faster. Notice there could be precision problems due to many successive floating point additions if `M` became much higher.

We are not sure whether or not those optimizations are substantial compared to the call to the mystery function, but it cannot do any harm anyways. Furthermore, since our class `SeqLaplaceInv` takes the function to invert as an argument, it might end up being useful to invert other functions less costly than our mystery function.

2.2 Parallel MPI algorithm

The program can be quite long when the size of the input increases. In order to accelerate what the machine does, we can use numerous processors in parallel. We call N_p the number of

processes. The program takes a vector of **double** as input, computes the algorithm for each data point and return a vector of **double** as output. We propose two different methods to parallelize the implementation:

- parallelizing the algorithm by rearranging the sequential algorithm,
- dividing the input into chunks, and running in parallel the sequential algorithm on each chunk.

In both implementations, every process has to treat a certain amount of data. If we call N_T the size of the total amount of data, each processor will treat a chunk of size $\lfloor \frac{N_T}{N_p} \rfloor$, and the "master" processor will treat the remaining $N_T - \lfloor \frac{N_T}{N_p} \rfloor$. It is not the most efficient way of handling the problem (even though in most cases the difference is negligible), but it is much easier to implement, because each processor except the master has the same amount of data to treat.

2.2.1 Parallel version of the algorithm

In the sequential version of the algorithm, we have numerous sums. We want to calculate these sums in parallel but the problem is that each term of each sum depends on previous sums and one or more other terms. In order to parallelize the algorithm we have to come up with steps that are independent from each other.

Maths behind the parallelization

The calculation of Fun is the part of the algorithm we want to parallelize.

$$sum = \frac{1}{2} \sum_{i=1}^{Ntr} (-1)^i fnRf(X, iH) \quad (1)$$

$$SU(1) = sum \quad (2)$$

$$SU(k+1) = SU(k) + (-1)^{Ntr+k} fnRf(X, (Ntr+k)H) \quad (3)$$

$$Fun = \frac{U}{Ctot} \sum_{i=1}^{M+1} \binom{M}{i-1} SU(i) \quad (4)$$

We first notice that we can easily calculate sum in parallel by calculating each term in parallel and then summing them. However, each $SU(i)$ needs $SU(i-1)$ to be calculated, so we must determine Fun differently to design the algorithm. We can rewrite (3) this way :

$$SU(k) = sum + \sum_{i=1}^{k-1} (-1)^{Ntr+i-1} fnRf(X, (Ntr+i-1)H) \quad (5)$$

In this case there are some values of $fnRf$ that are used by many SU , so that we lose time when we compute these expensive calculations more than once. From (4) and (5) we can rewrite a formulation for Fun :

$$Fun = \frac{U}{Ctot} \sum_{k=1}^{M+1} \binom{M}{k-1} \left(sum + \sum_{i=1}^{k-1} (-1)^{Ntr+i-1} fnRf(X, (Ntr+i-1)H) \right) \quad (6)$$

We can change (6) into :

$$S(1) = \text{sum} \quad (7)$$

$$S(k) = (-1)^{Ntr+k-1} fnRf(X, (Ntr + k - 1)H) \quad (8)$$

$$Fun = U \sum_{k=1}^{M+1} \left(S(k) \left(1 - \frac{\sum_{i=1}^{k-1} \binom{M}{i-1}}{Ctot} \right) \right) \quad (9)$$

Written this way, the terms of the sum in (9) are independent. So we can calculate each one of them in parallel.

Implementation

We use MPI to implement the parallel version of the algorithm. We set as "master" one of the processes. It will gather all the calculation results from the other processes and compute them to obtain the final result. For each value t of the input data, we run the parallel algorithm. The master is the only process that knows the input. Therefore the first step is to broadcast the t to all the processes. Then the algorithm is made of two major steps :

- the calculation of sum . This expression contains Ntr terms. Therefore each process calculates $\lfloor \frac{Ntr}{N_p} \rfloor$ terms, and then we perform a **Reduce** operation and the master gets the sum of all the terms. Concretely, after the calculation of one term by the processes, we perform a **Reduce** operation and the master saves a partial sum, which it updates at each step. If Ntr is not divisible by N_p , the master calculates the remaining terms of the sum.
- the calculation of Fun . It contains $M + 1$ terms, then in the same way as for sum the processes evaluates $\lfloor \frac{M+1}{N_p} \rfloor$ terms of the sum. Afterwards the MPI program does a **Reduce** operation and the master gets the total sum. It also deals with remaining terms if $M + 1$ is not divisible by N_p .

Finally, the master returns the value of Fun , and the operation is repeated for all the input.

2.2.2 Chunked method

The previous method performed a parallel calculation for each data point of the input. The chunked method does otherwise. It runs in parallel the sequential algorithm for different data points. We call N_{input} the size of the input. The chunked method also uses a process as master, and is divided into three parts:

- First the master, which is the only process knowing the input, allocates to all the processes the amount of data to be treated. Each process has to compute the algorithm on a chunk of $\lfloor \frac{N_{input}}{N_p} \rfloor$ points. Therefore the master scatters the input points to the processes.
- Then each process runs the sequential algorithm on the chunk of points that it has to process by calling the operator `()` of **SeqLaplaceInv**.
- Afterwards, we do a **Gather** to put all the results into a vector on the master. Then the masters deals with the remaining points if N_{input} is not divisible by N_p , by computing the sequential algorithm on them. Finally it returns the output vector.

2.3 Addition of OpenMP parallelization

Once we were sure that our algorithms using MPI were properly working, we decided to use OpenMP as well in order to increase their speed.

- The chunked method algorithm was very simple to parallelize using OpenMP. Since each process computes a chunk of inputs, we only needed to parallelize the `for` loops using the `#pragma omp parallel for` instruction. That way, the different threads compute separately each values and put them in the corresponding output vector.
- The parallel version of the sequential algorithm revealed itself to be extremely complicated to parallelize. Since we find MPI instructions in all the sections that could be parallelized using OpenMP, we were not able to correctly parallelize those sections. Here, we reach a limit in the use of both MPI and OpenMP at the same time.

2.4 Performance analysis

After implementing several solutions of the problem, we compare in this section their performances. The parameters that can be set are N , M and the size of the input. As previously mentioned, N and M influence the accuracy of the solution. Making them really large is not useful since there is a value above which the gain in precision is negligible (above all to make the clock readable). Therefore we set $N = 100$ and $M = 15$ in all our implementations, and we only study the influence of the input size on the performance.

The secret machine is a clock, therefore the input should be between 0 and 12. We then define the *step*, which corresponds to the interval between two points of the input. For example, for *step* equal to 0.1 there are 121 points in the input data. We obtain in Table 1 the performance of our implementations, defined as the time spent to complete the calculation, for several *steps*. MPI_1 corresponds to the parallelization of the algorithm, MPI_2 to the chunked method and MPI_{2O} to the chunked method with OpenMP. The results are given in seconds.

<i>step</i>	0.1	0.05	0.01	0.005	0.001
<i>Seq(s)</i>	16	33	149	288	1448
$MPI_1(s)$	3.42	6.84	32.44	64.70	328.57
$MPI_2(s)$	1.17	2.41	9.79	19.00	95.25
$MPI_{2O}(s)$	0.77	1.07	4.37	5.65	25.25

Table 1: Execution time of the different implementations measured for many sizes of input

In light of this results, we can make several observations. First we observe that the parallelization always gives a better performance than the sequential implementation, which reinforces the great importance of our secret mission.

Besides, the amount of time spent in the calculation is approximately proportional to the input size, except for the MPI_{2O} with small steps. It shows that some communication costs cannot be reduced, even with a small amount of data to treat. We also notice that the chunked method is better than the parallelization of the algorithm in every case.

It may be explained by the fact that the chunked method requires less communication between the processes, because they only exchange information twice, at the beginning and at the end of the calculation. Indeed, in the parallelization of the algorithm itself, the processes do communicate several times while processing each single point.

Finally, the use of OpenMP improves the efficiency of the program, allowing the program to be parallelize again on each process.

We compute in Table 2 the speed-up and the efficiency of the three parallel implementations for different step values. For MPI_1 and MPI_2 we parallelize the program on 16 computers, so there are 16 processors. But for the MPI_{2O} , we have 8 cores on each computer. So there are 128 processors. That explains the relatively bad results for the efficiency of MPI_{2O} .

<i>step</i>	0.1	0.05	0.01	0.005	0.001
$Seq(s)$	16	33	149	288	1448
$S_p(MPI_1)$	4.67	4.82	4.59	4.45	4.44
$E_p(MPI_1)$	0.29	0.30	0.29	0.28	0.28
$S_p(MPI_2)$	13.68	13.69	15.22	15.16	15.2
$E_p(MPI_2)$	0.86	0.86	0.95	0.95	0.95
$S_p(MPI_{2O})$	20.78	30.84	34.10	50.97	57.35
$E_p(MPI_{2O})$	0.16	0.24	0.27	0.40	0.45

Table 2: efficiency and speed-up of the different implementations

MPI_{2O} is obviously the fastest implementation although it is less efficient than MPI_2 .

3 Conclusion

Thanks to the parallel implementation, we manage to copy and surpass Nukehavistan technology. We achieve the same results in different and efficient ways.




Figure 4: Output of our program

But as we said before, we are French and French are all about beauty and class. The output of the program is very difficult to read, since it is very uncommon. Our spies want to act as quick as possible. We cannot conceive that they might lose some time reading time. Thus we tweak the output so that it is more readable.

This way, we are sure that our spies will have the best equipment while they are on-stage.




Figure 5: Left the clock; Right the basic output.




Figure 6: And we cannot say we worked off the clock

More about the program

Launching the program

The main part of the project is in the repository `implementation/`.

First you need to compile the programs. A complete *Makefile* is provided so that you only need the command `$ make`.

Then try `$./whattime` to get help about the syntax to use to launch the program.

Overall architecture

The general program contains 4 C++ classes:

- **LaplaceInv**: an abstract class that represents an inverse Laplace transform. It must be passed a function of the same signature as the mystery function `L`. The virtual operator `()` called on a `double` value or a `vector<double>` is meant to return its image by the inverse Laplace transform of that function. It is supposed to be inherited multiple times for the different types of implementation of this inverse.
- **SeqLaplaceInv** and **MPILaplaceInv**: they inherit from **LaplaceInv** and overload the operator `()`. The first one implements the sequential algorithm, whereas the second one is a wrapper that calls the other executable `mpicore`, which performs the computation in parallel using MPI.
- **Plotter**: it must be passed an object of type **LaplaceInv** at construction time and provides methods that call it and then plot the output using *GnuPlot*.

Its function `main()` is in *start.cpp* and instantiates one of those objects depending on the parameters given by the user.

Problems interfacing MPI with the rest of the program

Making a global program whose executable contains MPI code would be foolish since it would use several MPI processes each time or would require to be launched with `mpirun` in every case. Hence our decision to build a different executable to handle the MPI processing, and to launch it from the global program after forking.

The problem then is to make those two programs communicate. We need this since the wrapper must pass arguments to its child and retrieve output data from it.

The first option we considered was to make the data transit by name Unix pipes. After long hours of debug, such mechanisms appeared to be very painful to program in C++ (without using system calls directly), and might have raised problems since the distributed file system NFS is used in the lab environment.

Thus, we decided to use a less elegant (but easier to program) solution: writing the input and output data down to normal files. The only problem with this approach is that the `mpicore` process needs a way to wake the wrapper up when it is done processing the data and has finished writing the result to the output file. In our design, such a mechanism is performed by using TCP sockets: once it has launched the `mpicore` process, the wrapper opens a server socket on the loopback interface and waits for its child to connect to it, which triggers the reading of the output file, and then the plotting.



Nerds and Geeks Office
MI6 Building
85 Albert Embankment
London SE1

Telephone Bill

Number	Duration	Price
0700 0036	0h2m58s	£458.92
Total w/o VAT		£458.92
VAT 20%		£91.78
Total with VAT		£550.70